

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import seaborn as sns
        4 import matplotlib.pyplot as plt
        5 from sklearn import preprocessing, svm
        6 from sklearn.model_selection import train_test_split
        7 from sklearn.linear_model import LinearRegression
```

```
In [2]: 1 df=pd.read_csv(r"C:\Users\Welcome\Downloads\used_cars_data.csv")
        2 df
```

```
Out[2]:
```

	S.No.	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_T	
	0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	f
	1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	f
	2	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	f
	3	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	f
	4	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	Sec

	7248	7248	Volkswagen Vento Diesel Trendline	Hyderabad	2011	89411	Diesel	Manual	f
	7249	7249	Volkswagen Polo GT TSI	Mumbai	2015	59000	Petrol	Automatic	f
	7250	7250	Nissan Micra Diesel XV	Kolkata	2012	28000	Diesel	Manual	f
	7251	7251	Volkswagen Polo GT TSI	Pune	2013	52262	Petrol	Automatic	T
	7252	7252	Mercedes-Benz E-Class 2009-2013 E 220 CDI Avan...	Kochi	2014	72443	Diesel	Automatic	f

7253 rows × 14 columns



In [3]: 1 df.head()

Out[3]:

	S.No.	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type
0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	First
1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	First
2	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	First
3	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	First
4	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	Second

In [4]: 1 df.tail()

Out[4]:

	S.No.	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Ty
7248	7248	Volkswagen Vento Diesel Trendline	Hyderabad	2011	89411	Diesel	Manual	F
7249	7249	Volkswagen Polo GT TSI	Mumbai	2015	59000	Petrol	Automatic	F
7250	7250	Nissan Micra Diesel XV	Kolkata	2012	28000	Diesel	Manual	F
7251	7251	Volkswagen Polo GT TSI	Pune	2013	52262	Petrol	Automatic	Th
7252	7252	Mercedes-Benz E-Class 2009-2013 E 220 CDI Avan...	Kochi	2014	72443	Diesel	Automatic	F

In [5]: 1 df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7253 entries, 0 to 7252
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   S.No.                 7253 non-null  int64  
 1   Name                  7253 non-null  object  
 2   Location              7253 non-null  object  
 3   Year                  7253 non-null  int64  
 4   Kilometers_Driven     7253 non-null  int64  
 5   Fuel_Type             7253 non-null  object  
 6   Transmission          7253 non-null  object  
 7   Owner_Type            7253 non-null  object  
 8   Mileage               7251 non-null  object  
 9   Engine                7207 non-null  object  
10  Power                 7207 non-null  object  
11  Seats                 7200 non-null  float64 
12  New_Price             1006 non-null  object  
13  Price                 6019 non-null  float64 
dtypes: float64(2), int64(3), object(9)
memory usage: 793.4+ KB

```

In [6]: 1 df.describe()

Out[6]:

	S.No.	Year	Kilometers_Driven	Seats	Price
count	7253.000000	7253.000000	7.253000e+03	7200.000000	6019.000000
mean	3626.000000	2013.365366	5.869906e+04	5.279722	9.479468
std	2093.905084	3.254421	8.442772e+04	0.811660	11.187917
min	0.000000	1996.000000	1.710000e+02	0.000000	0.440000
25%	1813.000000	2011.000000	3.400000e+04	5.000000	3.500000
50%	3626.000000	2014.000000	5.341600e+04	5.000000	5.640000
75%	5439.000000	2016.000000	7.300000e+04	5.000000	9.950000
max	7252.000000	2019.000000	6.500000e+06	10.000000	160.000000

In [7]: 1 df.shape

Out[7]: (7253, 14)

```
In [8]: 1 df.isna().any()
```

```
Out[8]: S.No.                False
        Name                False
        Location            False
        Year                False
        Kilometers_Driven    False
        Fuel_Type            False
        Transmission         False
        Owner_Type           False
        Mileage              True
        Engine               True
        Power                True
        Seats                True
        New_Price            True
        Price                True
        dtype: bool
```

```
In [9]: 1 df.isnull().sum()
```

```
Out[9]: S.No.                0
        Name                0
        Location            0
        Year                0
        Kilometers_Driven    0
        Fuel_Type            0
        Transmission         0
        Owner_Type           0
        Mileage              2
        Engine               46
        Power                46
        Seats                53
        New_Price            6247
        Price                1234
        dtype: int64
```

```
In [11]: 1 df.fillna(value=0,inplace=True)
```

```
In [12]: 1 df.isnull().sum()
```

```
Out[12]: S.No.                0
        Name                0
        Location            0
        Year                0
        Kilometers_Driven    0
        Fuel_Type            0
        Transmission         0
        Owner_Type           0
        Mileage              0
        Engine               0
        Power                0
        Seats                0
        New_Price            0
        Price                0
        dtype: int64
```

```
In [13]: 1 df=df[['Year','Price']]
        2 df.columns=['Yr','Pr']
```


```
In [15]: 1 df.head(10)
```

```
Out[15]:
```

	Yr	Pr
0	2010	1.75
1	2015	12.50
2	2011	4.50
3	2012	6.00
4	2013	17.74
5	2012	2.35
6	2013	3.50
7	2016	17.50
8	2013	5.20
9	2012	1.95

```
In [17]: 1 sns.lmplot(x="Yr",y="Pr",data=df,order=2,ci=None)
```

```
Out[17]: <seaborn.axisgrid.FacetGrid at 0x20afd13ccd0>
```



In [18]: 1 df.describe()

Out[18]:

	Yr	Pr
count	7253.000000	7253.000000
mean	2013.365366	7.866665
std	3.254421	10.796286
min	1996.000000	0.000000
25%	2011.000000	2.290000
50%	2014.000000	4.650000
75%	2016.000000	8.400000
max	2019.000000	160.000000

In [19]: 1 df.fillna(method='ffill',inplace=True)

C:\Users\Welcome\AppData\Local\Temp\ipykernel_15892\4116506308.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

df.fillna(method='ffill',inplace=True)

In [20]: 1 x=np.array(df['Yr']).reshape(-1,1)
2 y=np.array(df['Pr']).reshape(-1,1)
3

In [21]: 1 df.dropna(inplace=True)

C:\Users\Welcome\AppData\Local\Temp\ipykernel_15892\1379821321.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

df.dropna(inplace=True)

In [22]: 1 X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
2 reg=LinearRegression()
3 reg.fit(X_train,y_train)
4 print(reg.score(X_test,y_test))


0.060339820443730874

```
In [23]: 1 y_pred=reg.predict(X_test)
2 plt.scatter(X_test,y_test,color='b')
3 plt.plot(X_test,y_pred,color='k')
4 plt.show()
```




```
In [24]: 1 df500=df[:][:500]  
2 sns.lmplot(x="Yr",y="Pr",data=df500,order=1,ci=None)
```

Out[24]: <seaborn.axisgrid.FacetGrid at 0x20aff25e230>




```
In [25]: 1 df500.fillna(method='ffill',inplace=True)
2 X=np.array(df500['Yr']).reshape(-1,1)
3 y=np.array(df500['Pr']).reshape(-1,1)
4 df500.dropna(inplace=True)
5 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
6 regr=LinearRegression()
7 regr.fit(X_train,y_train)
8 print("Regression:",regr.score(X_test,y_test))
9 y_pred=regr.predict(X_test)
10 plt.scatter(X_test,y_test,color='b')
11 plt.plot(X_test,y_pred,color='k')
12 plt.show()
13
```

Regression: 0.179288978213473



```
In [26]: 1 from sklearn.linear_model import LinearRegression
2 from sklearn.metrics import r2_score
3 model=LinearRegression()
4 model.fit(X_train,y_train)
5 y_pred=model.predict(X_test)
6 r2=r2_score(y_test,y_pred)
7 print("R2 score:",r2)
```

R2 score: 0.179288978213473

